



Machine Learning *Lecture 7*

David Meyer
Pascal Plank

Lecture 7 – Unsupervised Learning

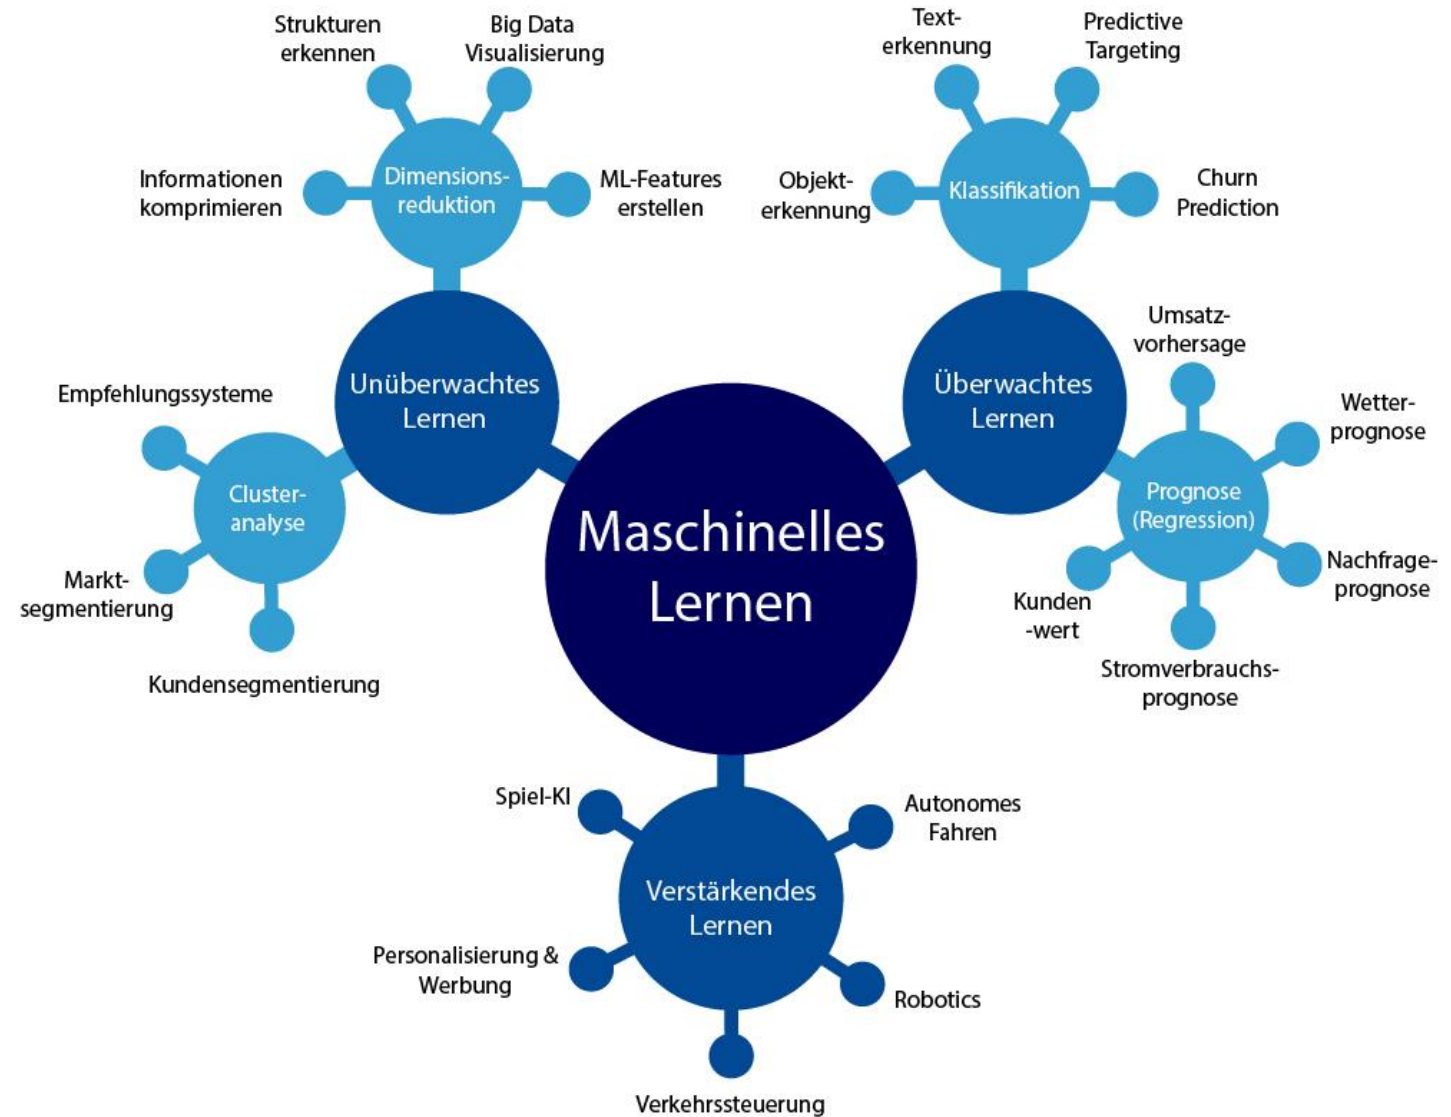
- **Intro Unsupervised Learning**
- **Clustering**
- **K-means clustering**
- **Hierarchical clustering**
- **DBSCAN**
- **Association rules**

Lecture 7 – Unsupervised Learning

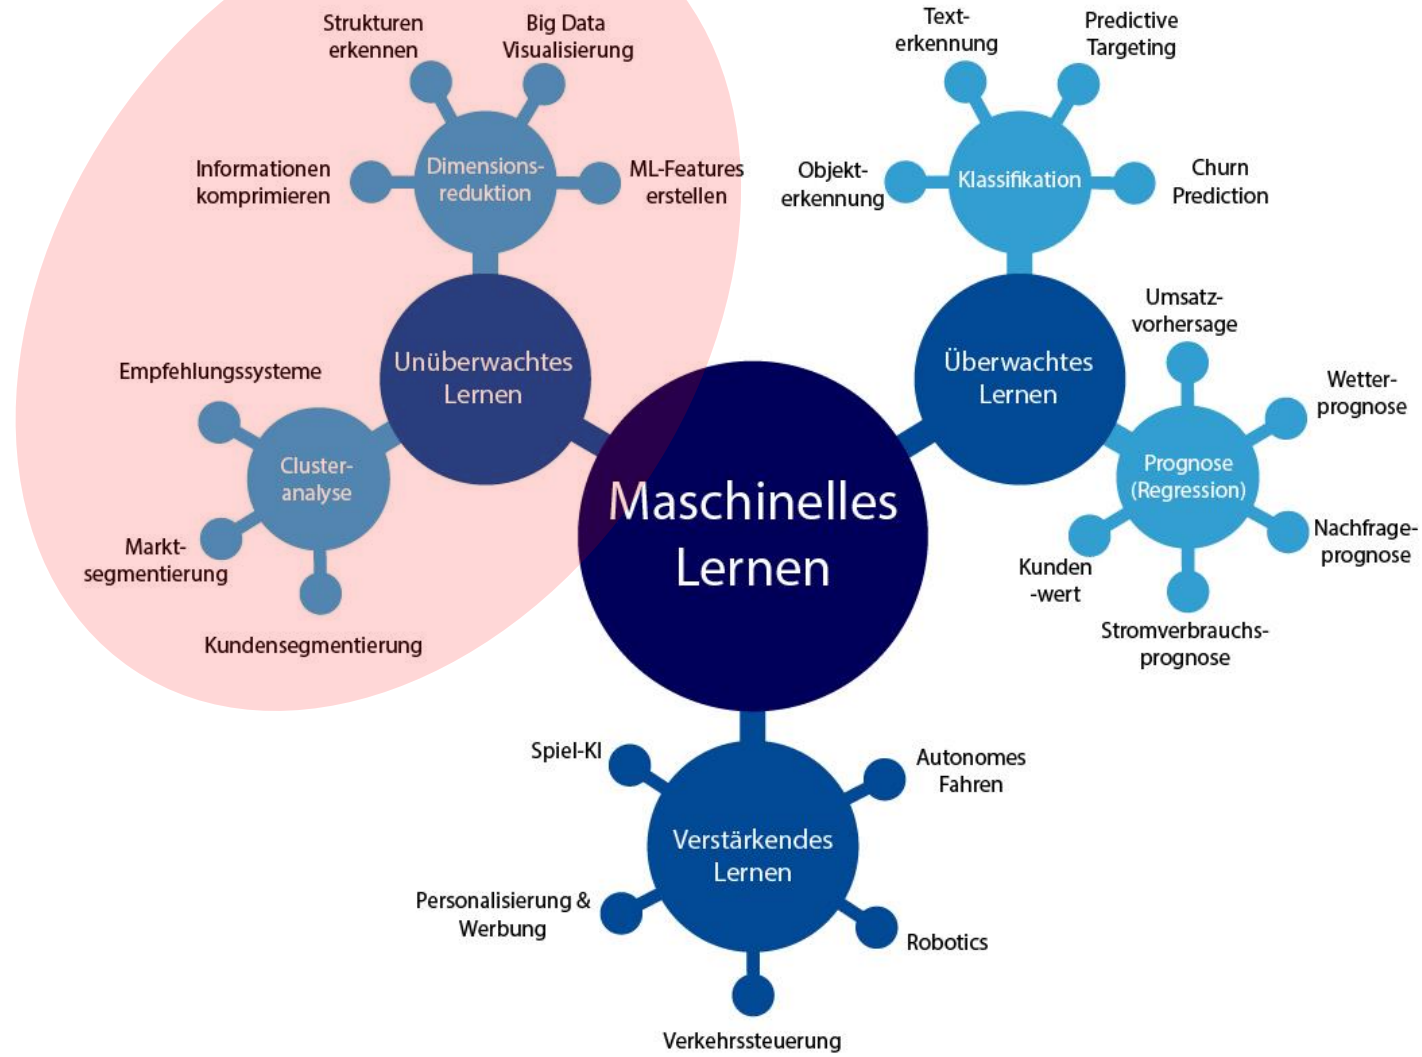
- **Intro Unsupervised Learning**
- **Clustering**
- **K-means clustering**
- **Hierarchical clustering**
- **DBSCAN**
- **Association rules**

Machine Learning Types

- **Supervised Learning:**
 - labelled data
 - Direct feedback
 - Predict an outcome/future, forecasting
 - E. g. predict customer churn
- **Unsupervised Learning:**
 - No labels
 - No feedback
 - Finding hidden structures
 - E. g. cluster customers
- **Reinforcement Learning:**
 - Decision process
 - Reward system
 - Learn a series of actions
 - E. g. playing Go




Machine Learning Types



Why unsupervised learning?

- Getting enough ground truth data (a.k.a. labelled data) is expensive
- Automatic collection would be nice
- Features can update and the labels should update too
→ e. g. color of fruit over a season; age/growth rate/weight of a person etc.
- Use cases are quite diverse, some examples:
 - Data analysis in science (e. g. grouping of genes, proteins, astronomical objects)
 - Marketing (e. g. product analysis, shopping basket analysis, grouping of customers)
 - Databases (e. g. finding similar documents)



Lecture 7 – Unsupervised Learning

- Intro Unsupervised Learning
- **Clustering**
- K-means clustering
- Hierarchical clustering
- DBSCAN
- Association rules

Clustering

- Unsupervised learning method to cluster data into groups
- **Goals:**
 - Clusters should be homogeneous (samples in a single cluster are similar)
 - Clusters themselves should be very different
- **Definition:**

m-clustering of $X = \{x_1, x_2, \dots, x_N\}$ is a subdivision of X into m cluster C_1, \dots, C_m such that:

- $C_i \neq \emptyset, i = 1, \dots, m$
- $\bigcup_{i=1}^m C_i = X$
- $C_i \cap C_j = \emptyset, i \neq j; i, j = 1, \dots, m$

The vectors x_i of cluster C_i are more similar to each other than they are to the vectors of the other clusters. This is called **„hard“ or „crisp“ clustering**.

Clustering

- **Definition (Fuzzy Clustering):**

Fuzzy-clustering of $X = \{x_1, x_2, \dots, x_N\}$ into m cluster C_1, \dots, C_m is defined by m membership functions u_j with:


- $u_i: X \rightarrow [0,1], j = 1, \dots, m$ and
- $\sum_{j=1}^m u_i(x_i) = 1, i = 1, \dots, N$

Each vector x_i is part of several clusters at the same time, the value of this membership is quantified by u_j . Values near 1 indicate a high degree of membership to a cluster.

Clustering

- There are many possible solutions depending on the scale

[Quelle: Theodoris et al., 2009]




Clustering

- For clustering to work, you'll need:
 - Features and some pre-processing
 - Distance metric
 - Cluster criterion (grouping criterion)
 - Clustering algorithm
 - Validation and interpretation of results
- Depending on the selection of the above points, you'll get very different clusterings

Clustering

[Quelle: http://scikit-learn.org/stable/auto_examples/cluster/plot_cluster_comparison.html]



Clustering – Similarity metrics

- **Euclidean norm**
- **dot-product** (scalar product) of two vectors x, y :

$$s_{scalar}(x, y) = x^T y = \sum_{i=1}^l x_i y_i$$


- **Pearson's correlation coefficient:**

$$s_{pearson}(x, y) = \frac{x_d^T y_d}{\|x_d\| \|y_d\|},$$

with $x_d[x_1 - \bar{x}, \dots, x_l - \bar{x}]^T$ and $y_d[y_1 - \bar{y}, \dots, y_l - \bar{y}]^T$

Clustering – Similarity metrics

- **Distance between a vector and a subset of X**
 - needed for assessing whether a vector should be assigned to a cluster C
- Several options exist:
 - Minimal dissimilarity: $d(x,C) = \min(d(x,y))$ for all y in C
 - Maximal similarity: $s(x,C) = \max(s(x,y))$ for all y in C
 - Average (dis)similarity or distance to a **centroid** average cluster coordinates)
- Depending on the cluster type, a different distance calculation might be reasonable



kompakt



länglich



kreisförmig, elliptisch

[Quelle: Theodoris et al., 2009]

Lecture 7 – Unsupervised Learning

- Intro Unsupervised Learning
- Clustering
- **K-means clustering**
- Hierarchical clustering
- DBSCAN
- Association rules

K-means clustering

- K-means is a **centroid-based clustering method. It is partitional**
- **Goal:** find the k cluster centers and assign the objects to the nearest cluster center (centroids), such that the squared distances from the cluster are minimized
- **NP-hard** problem → we search for an approximate solution
- Number of clusters k needs to be specified (hyperparameter)
- Conceptually close to k-nearest neighbor classification algorithm
- Some variations:
 - k-median (more robust against outliers)
 - k-medoids (centroids need to be a data sample)
 - k-means++ (choosing initial centers less randomly)
 - fuzzy c-means (allowing fuzzy cluster assignment)

K-means clustering


- **The algorithm consists of the following steps:**
 1. Choose number of clusters k and initial centroids (usually randomly)
 2. Assignment step: assign samples to the cluster with the closest (Euclidean distance) centroid
 3. Update step: Recalculate centroids (means) with the assignments of step 2
 4. Repeat steps 2 and 3 until there is no significant improvement anymore
- The algorithm is not guaranteed to find the global optimum
- It is based on spherical clusters that are separable and approximately the same size so that the mean converges towards the cluster center

K-means clustering

- To assess „significant improvement“ we need a criterion, e. g., sum of squared distances of samples to their centroids
- Since we use distance metrics **standardized variables** are strongly recommended!
- Initial centroids are critical to the overall result
 - repeat algorithm several times with random initial centroids
 - choose the best run
- The resulting clusters can be used as an additional categorical variables for further statistical methods, e. g., clustering customers and running analysis on customer clusters.

K-means clustering – Choosing k

- K is very important to the final result
- To find a good k, run the algorithm several times with different k and plot the errors



- Look out for a plateauing point or a sharp change in error trend that indicates that improvement with higher k is limited after this point.

K-means clustering – Evaluating results with silhouettes

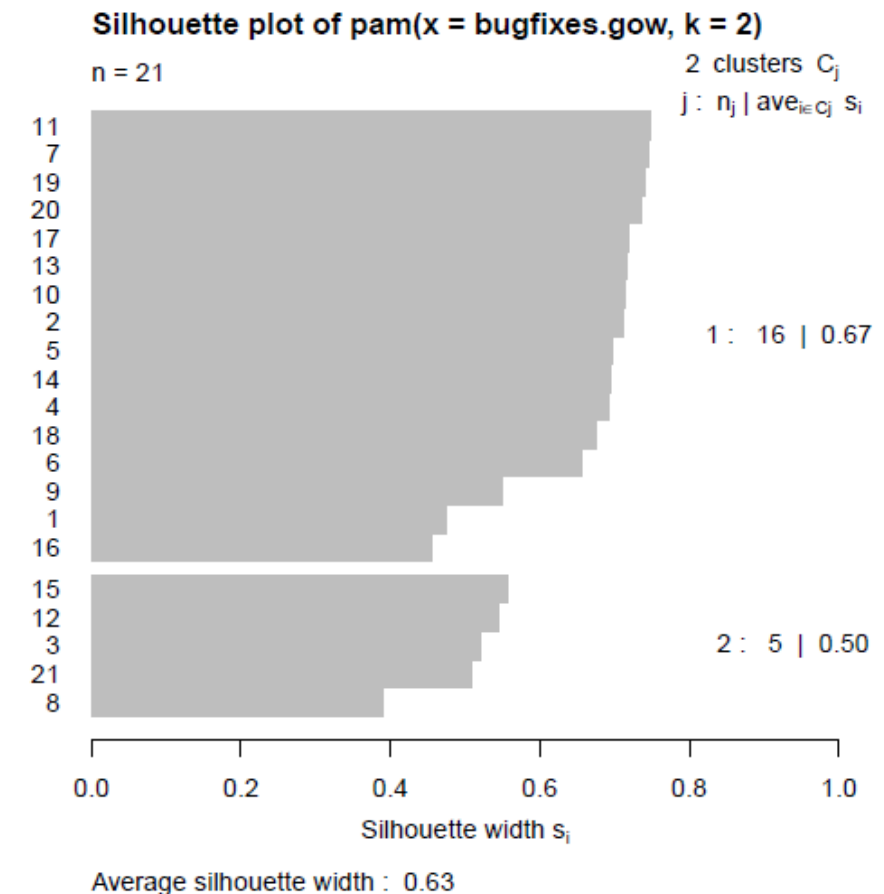
- How good the samples are clustered can be assessed with a silhouette plots
- Silhouettes can be calculated for the overall solution or individual clusters

- Silhouette values and their approx. meaning:

- 1.0 – 0.71: well clustered samples, clustering found strong structures
- 0.7 – 0.51: medium and somewhat acceptable structures were found
- 0.5 – 0.26: weak structures
- <0.25 : no substantial structures were found, and samples lie between clusters

- Example for $n=21$ samples in 2 clusters:

- Clusters are medium (0.5) or strong clusters (0.67)
- Average silhouette width is acceptable (0.63)
- Individual silhouettes are >0 , i. e., no samples are in wrong cluster or close to a border




Class Assignment

- **Assignment:**


Use **K=2** with points **A** and **C** as initial cluster means and perform the **k-means clustering algorithm** for **two iterations** (Euclidean norm $d = \sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2}$)


	i	x_1	x_2
\bar{X}_1^0	A	1	1
	B	1	0
\bar{X}_2^0	C	0	2
	D	2	4
	E	3	5



Class Assignment – Solution 1/7

- Step 1 (Iteration 1): Calculate the distances between each of the cluster means and all other points

\bar{X}_1^0 

\bar{X}_2^0 

i	X_1	X_2
A	1	1
B	1	0
C	0	2
D	2	4
E	3	5

i	1	2
A	0	1.4
B	1	2.2
C	1.4	0
D	3.2	2.8
E	4.5	4.2

Class Assignment – Solution 2/7

- Step 2 (Iteration 1): Assign each point to the cluster with the lowest distance. Use this new cluster assignment to calculate the new cluster means


i	1	2	Cluster
A	0	1.4	1
B	1	2.2	1
C	1.4	0	2
D	3.2	2.8	2
E	4.5	4.2	2

i	X ₁	X ₂
A	1	1
B	1	0
C	0	2
D	2	4
E	3	5

 \bar{X}_1^1  \bar{X}_2^1

Class Assignment – Solution 3/7

- Iteration 1 final: Updated clusters in the graph:



● $\bar{X}_1^1 = (1, 0.5)$

● $\bar{X}_2^1 = (1.7, 3.7)$

Class Assignment – Solution 4/7

- Step 3 (Iteration 2): Recalculate the distances of points to the new cluster means

i	X ₁	X ₂
A	1	1
B	1	0
C	0	2
D	2	4
E	3	5

● $\bar{X}_1^1 = (1, 0.5)$

● $\bar{X}_2^1 = (1.7, 3.7)$

i	1	2
A	0.5	2.7
B	0.5	3.7
C	1.8	2.4
D	3.6	0.5
E	4.9	1.9

Class Assignment – Solution 5/7

- Step 4 (Iteration 2): Assign each point to the cluster with the lowest distance.

i	1	2	Cluster
A	0.5	2.7	1
B	0.5	3.7	1
C	1.8	2.4	1
D	3.6	0.5	2
E	4.9	1.9	2

i	X_1	X_2
A	1	1
B	1	0
C	0	2
D	2	4
E	3	5

 \bar{X}_1^2

 \bar{X}_2^2

Class Assignment – Solution 6/7

- Step 5 (Iteration 2): Assign each point to the cluster with the lowest distance.

i	1	2	Cluster
A	0.5	2.7	1
B	0.5	3.7	1
C	1.8	2.4	1
D	3.6	0.5	2
E	4.9	1.9	2

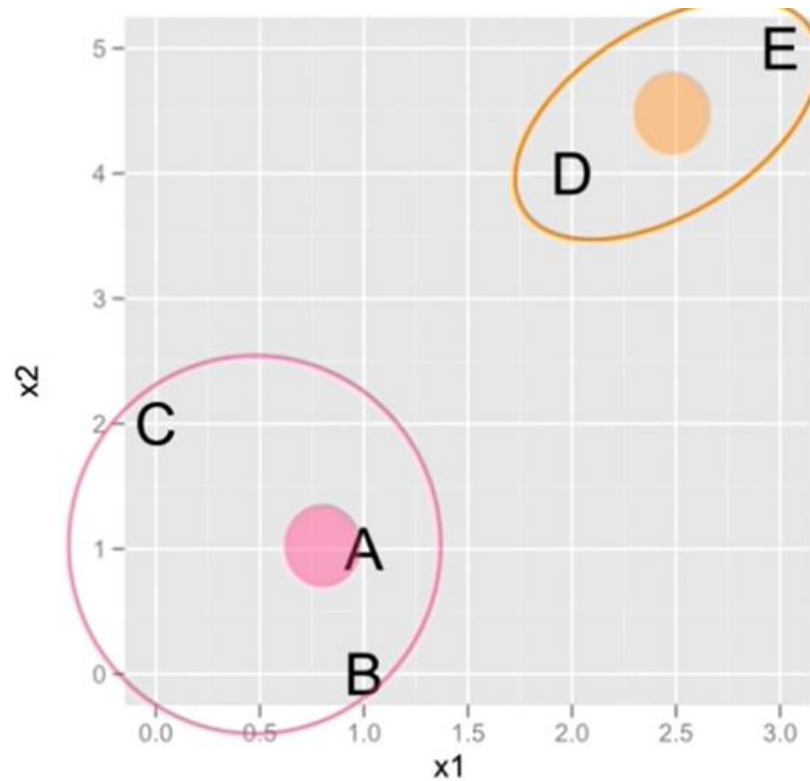
i	X ₁	X ₂
A	1	1
B	1	0
C	0	2
D	2	4
E	3	5

 \bar{X}_1^2

 \bar{X}_2^2

Class Assignment – Solution 7/7

- Step 6 (Iteration 2): Recalculate cluster means
→ Algorithm has already converged 😊



● $\bar{X}_1^2 = (0.7, 1)$

● $\bar{X}_2^2 = (2.5, 4.5)$

Lecture 7 – Unsupervised Learning

- Intro Unsupervised Learning
- Clustering
- K-means clustering
- **Hierarchical clustering**
- DBSCAN
- Association rules

Hierarchical clustering

- **Two types:**
 - Agglomerative
 - Divisive
- In addition to homogeneous clusters, hierarchical clustering results in a hierarchical structure of the data (a **dendrogram**)

Clustering-Hierarchie


$\{\{x_1\}, \{x_2\}, \{x_3\}, \{x_4\}, \{x_5\}\}$

$\{\{x_1, x_2\}, \{x_3\}, \{x_4\}, \{x_5\}\}$

$\{\{x_1, x_2\}, \{x_3\}, \{x_4, x_5\}\}$

$\{\{x_1, x_2\}, \{x_3, x_4, x_5\}\}$

$\{\{x_1, x_2, x_3, x_4, x_5\}\}$




Dendrogram



[Quelle: Theodoris et al., 2009]

Hierarchical clustering – agglomerative approach

- Start with each sample in a separate cluster and then group them into larger and larger clusters until there is only one cluster left
- Grouping criterion determines how this agglomeration works:
 - Single Linkage Clustering: smallest distance between samples in clusters
 - Complete Linkage Clustering: largest distance between samples in clusters
 - Average Linkage Clustering: mean distance between samples in clusters




Hierarchical clustering – agglomerative approach



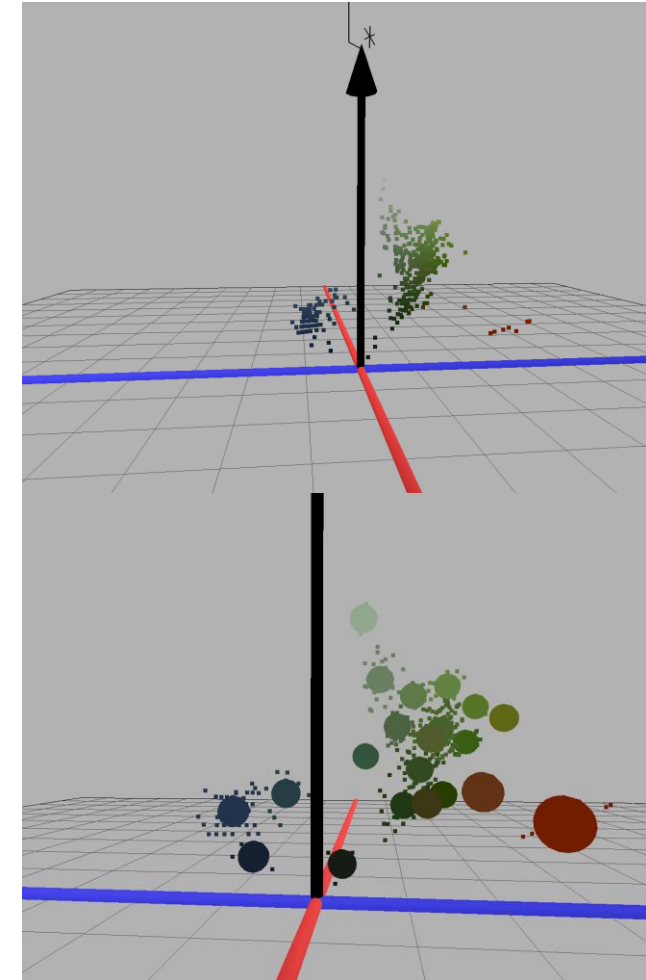
Hierarchical clustering – divisive approach

- Opposite of agglomerative approach: start with all samples in one cluster and recursively split until each sample is in its own cluster
- Naïve method: Search for best subdivision of a set into two clusters out of the $2^{N-1}-1$ possibilities based on some criterion.



Hierarchical clustering – Remarks

- There is no way to recover from a “bad grouping” (or splitting)
- To get k clusters just stop early or cut off the appropriate branches in the dendrogram
- Example: reduce color palette to 500 colors (or even fewer)




Lecture 7 – Unsupervised Learning


- Intro Unsupervised Learning
- Clustering
- K-means clustering
- Hierarchical clustering
- **DBSCAN**
- Association rules

DBSCAN



- DBSCAN ... “*Density-Based Spatial Clustering of Applications with Noise*”
- Works well even with “nested clusters” (also in higher dimensions with the right distance metric)
+ reduced risk of *single-link clusters*
- Sequential algorithm that also marks outliers
- **Hyperparameters:** distance d , min. neighbors n





DBSCAN – Algorithm




1.) Starting point: raw unlabelled points





2.) Count the number of points within **distance d** (orange, hyperparameter) for each point




3.) Label points as „**core**“ points (red) if they have more than **n neighbors** (hyperparameter, e. g. 4) within **distance d** , otherwise as „**non-core**“ points (black)



4.) Each point is either **core** or **non-core**




5.) Randomly pick a **core point** and assign it to **cluster 1**. Add all core points within **distance d** to that point to the cluster.




6.) Repeat step 5 for all added core points recursively until there are no more core points to add


DBSCAN – Algorithm




7.) Add all **non-core points** within distance d to **cluster 1** (only once, not recursively!)



8.) That creates our **first cluster (green)**



9.) Like in step 5, pick a new **non-clustered core point as a seed for cluster 2 (blue)** and recursively add all core points closer than distance d



10.) Like in steps 7 & 8, add the **non-core points** and we now have our **full cluster 2**

11.) If there are more non-clustered core points left, repeat the steps until there are no more unclustered core points left.

All remaining non-core points are marked as **outliers**.

DBSCAN – Remarks

1. No initial number of clusters necessary
2. Mostly deterministic except for some non-core boundary points and cluster order
3. Run time complexity $O(n^2)$ but some tradeoffs can be made by using efficient memory structures (e. g. R-tree variations) and pre-calculations
4. Robust to outliers (we even get a list of them)
5. Clusters should have similar densities for the algorithm to work properly (otherwise the min neighbor parameter only fits selected clusters well)

Lecture 7 – Unsupervised Learning

- Intro Unsupervised Learning
- Clustering
- K-means clustering
- Hierarchical clustering
- DBSCAN
- **Association rules**

Association rules

- Association Rule Mining is a rule-based machine learning method that helps to uncover meaningful correlations between different products according to their co-occurrence in a data set
- in sales/marketing also called “Market Basket Analysis”
- Association rules represent relationships between set of elements in transactions (not an individual's preference)
- “Customers also bought xxx” → Recommendation
- “Frequently bought together items” → Association
- Market Basket Analysis is very simple:
→ Basic idea: check likelihood of different elements occurring together
- We use the Apriori-Algorithm
- Association rules use:
 - **Support**
 - **Confidence**
 - **Lift**

Wird oft zusammen gekauft



Gesamtpreis: **167,77 €**

Alle drei in den Einkaufswagen

Example Use Cases:

- A and B can be placed together so that when a customer buys one of the product he doesn't have to go far away to buy the other product.
- People who buy one of the products can be targeted through an advertisement campaign to buy the other.
- Collective discounts can be offered on these products if the customer buys both of them.
- Both A and B can be packaged together.

Association rules

Rules need to meet a **user-specified minimum support and minimum confidence**

The diagram illustrates the calculation of three metrics for an association rule $Rule: X \Rightarrow Y$. Three blue arrows originate from the rule and point to the respective formulas:

- Support** is calculated as $Support = \frac{freq(X, Y)}{N}$.
- Confidence** is calculated as $Confidence = \frac{freq(X, Y)}{freq(X)}$.
- Lift** is calculated as $Lift = \frac{Support}{Supp(X) \times Supp(Y)}$.

- **Support** ... (general) probability that X occurs over all transactions
- **Confidence** ... conditional probability of Y when X occurred/is present
- **Lift** ... how likely an item is purchased when another item is purchased, while controlling for how popular both items are
(higher value == higher association; >1 ... "association is present")


Association rules – Apriori algorithm

Apriori principle: if an itemset is infrequent, then all its supersets must also be infrequent
→ use this for effective pruning of possible combinations to reduce the required computational effort

Steps:

1. **Hyperparameter Setting:** set a minimum value for **support** and **confidence**. This means that we are only interested in finding rules for the items that have certain default existence (i. e. support) and have a minimum value for co-occurrence with other items (i. e. confidence)
2. **Frequent Itemset Generation:** extract all the itemsets having higher value of support than minimum threshold
 - a) Start with itemset containing only 1 element, e. g., {bread}, {butter}
 - b) Calculate the support for itemsets and only keep those itemsets that reach a specified minimum support threshold (*on-the-fly-pruning – see next slide for animation*)
 - c) Generate all possible (non-empty) itemset combinations
 - d) Repeat step b and c until there are no more new valid itemset combinations
3. **Strong Rule Generation:** generate all possible rules from the itemsets and only select those with confidence values higher than the minimum threshold
4. **Order list of association rules:** order the rules by descending order of Lift

Association rules – Apriori algorithm animation (pruning via Support)

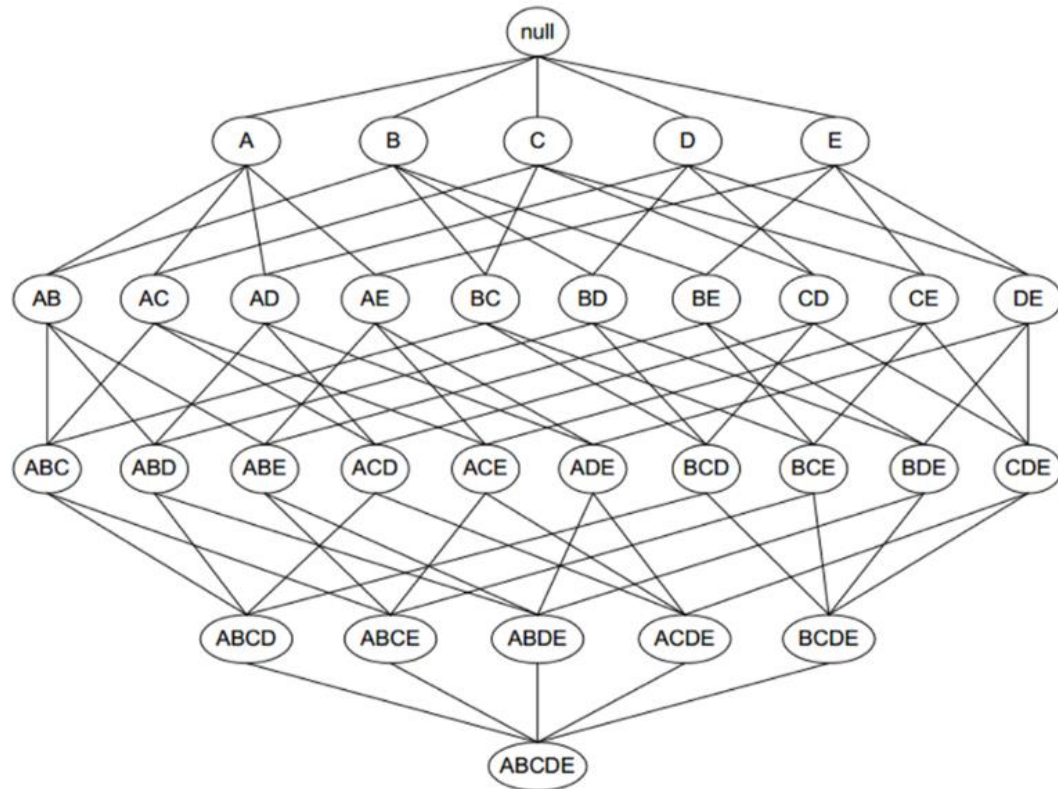


<https://www.kdnuggets.com/2016/04/association-rules-apriori-algorithm-tutorial.html/2>

Association rules – Apriori algorithm animation

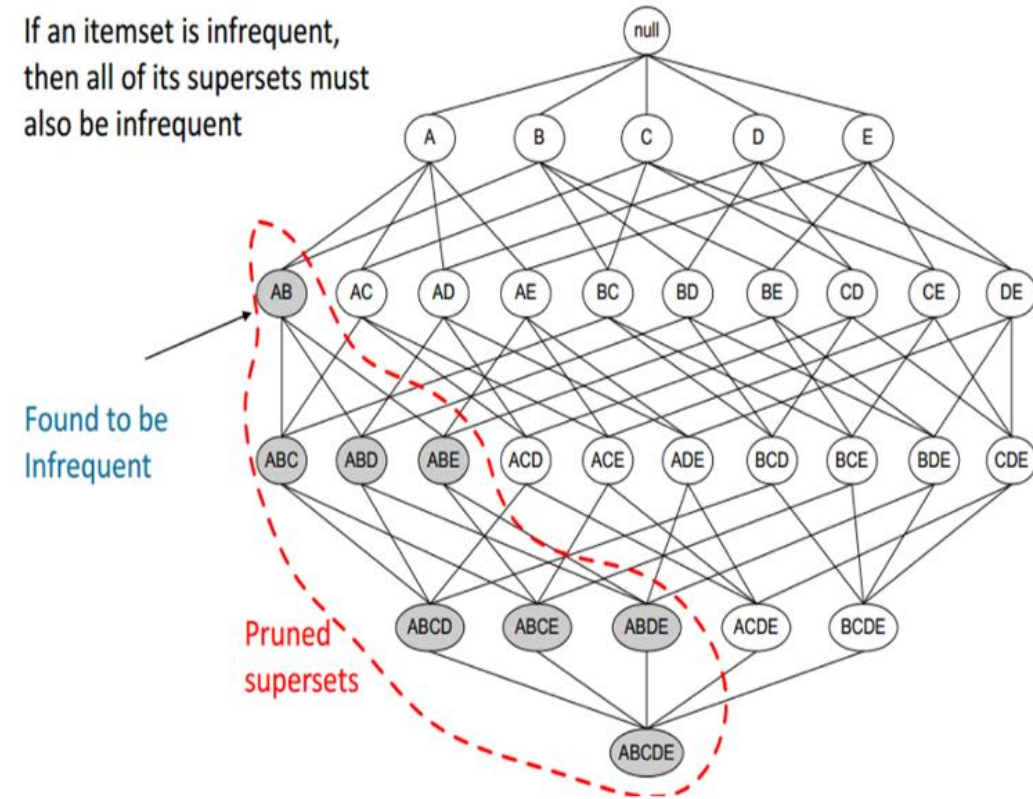
(pruning via Support)

The combinations of 5 items



The Apriori Algorithm

If an itemset is infrequent,
then all of its supersets must
also be infrequent



<https://www.kdnuggets.com/2019/12/market-basket-analysis.html>

Apriori algorithm example

For this example we will use a minimum support of 2 and a minimum confidence of 50%

TID	items
T1	I1, I2 , I5
T2	I2,I4
T3	I2,I3
T4	I1,I2,I4
T5	I1,I3
T6	I2,I3
T7	I1,I3
T8	I1,I2,I3,I5
T9	I1,I2,I3

Apriori algorithm example

Frequent Itemset Generation:

1. Calculate support values and only keep those with values larger than 2 (itemset size 1)
2. Combine sets and check again for the support value (itemset size 2)
3. Combine those sets and check again for the support value (itemset size 3)

(→ Remember: we use sets, so we do not care about the order of items in an itemset)

TID	items
T1	I1, I2, I5
T2	I2, I4
T3	I2, I3
T4	I1, I2, I4
T5	I1, I3
T6	I2, I3
T7	I1, I3
T8	I1, I2, I3, I5
T9	I1, I2, I3

original dataset

Itemset	sup_count
I1	6
I2	7
I3	6
I4	2
I5	2

Step 1 (Level 1)

Itemset	sup_count
I1,I2	4
I1,I3	4
I1,I4	1
I1,I5	2
I2,I3	4
I2,I4	2
I2,I5	2
I3,I4	0
I3,I5	1
I4,I5	0

Step 2 (Level 2) with filtering for min support of 2

Itemset	sup_count
I1,I2	4
I1,I3	4
I1,I5	2
I2,I3	4
I2,I4	2
I2,I5	2
I2,I5	2

Itemset	sup_count
I1,I2,I3	2
I1,I2,I5	2

Step 3 (Level 3)
with filtering for min
support of 2

Apriori algorithm example

Strong Rule Generation:

For each frequent itemset:

1. Generate all combinations of possible rules with the items in your itemset
2. Calculate the Confidence values
3. Only keep those with Confidence > min Confidence

TID	items
T1	I1, I2 , I5
T2	I2,I4
T3	I2,I3
T4	I1,I2,I4
T5	I1,I3
T6	I2,I3
T7	I1,I3
T8	I1,I2,I3,I5
T9	I1,I2,I3

original dataset

Example for {I1, I2, I3}:

- $[I1 \wedge I2] \Rightarrow [I3] \text{ //confidence} = \frac{\text{sup}(I1 \wedge I2 \wedge I3)}{\text{sup}(I1 \wedge I2)} = \frac{2}{4} * 100 = 50\%$
- $[I1 \wedge I3] \Rightarrow [I2] \text{ //confidence} = \frac{\text{sup}(I1 \wedge I2 \wedge I3)}{\text{sup}(I1 \wedge I3)} = \frac{2}{4} * 100 = 50\%$
- $[I2 \wedge I3] \Rightarrow [I1] \text{ //confidence} = \frac{\text{sup}(I1 \wedge I2 \wedge I3)}{\text{sup}(I2 \wedge I3)} = \frac{2}{4} * 100 = 50\%$
- $[I1] \Rightarrow [I2 \wedge I3] \text{ //confidence} = \frac{\text{sup}(I1 \wedge I2 \wedge I3)}{\text{sup}(I1)} = \frac{2}{6} * 100 = 33\%$
- $[I2] \Rightarrow [I1 \wedge I3] \text{ //confidence} = \frac{\text{sup}(I1 \wedge I2 \wedge I3)}{\text{sup}(I2)} = \frac{2}{7} * 100 = 28\%$
- $[I3] \Rightarrow [I1 \wedge I2] \text{ //confidence} = \frac{\text{sup}(I1 \wedge I2 \wedge I3)}{\text{sup}(I3)} = \frac{2}{6} * 100 = 33\%$

Example: A confidence of 50% means that 50% of the customers, who purchased I1 and I2 also bought I3.

Apriori algorithm example

Lift:

- We now have rules of frequent itemsets where we are confident that there is a correlation
 - BUT: some items are just bought very often and might create fluke rules
- use a final ordering of rule via the Lift value to sort that out

Example:

Transaction	Support	Confidence	Lift
Canned Beer → Soda	1%	20%	1.0
Canned Beer → Berries	0.1%	1%	0.3
Canned Beer → Male Cosmetics	0.1%	1%	2.6

- Lift = 1: no actual association (because beer and soda are just bought quite often)
- Lift < 1: negative association (when somebody buys beer they are usually not buying berries)
- Lift > 1: positive association (when somebody buys beer they buy male cosmetics – although the effect is small with a confidence of 1%)

Association rules – Apriori algorithm

- Drawbacks:
 - Computationally expensive
→ time and space complexity is exponential $O(2^{|D|})$ with $|D|$ being the horizontal width (= total number of items)
 - Spurious associations: some rules may seem logical
- Common improvements:
 - Hash-based itemset counting
 - Transaction reduction: A transaction that does not contain any frequent k-itemsets useless in subsequent scans
 - Partitioning: An itemset that is potentially frequent must be frequent in at least one of the partitions
 - Sampling: Mining on a subset of given data, lower support threshold + a method to determine the completeness
 - Dynamic itemset counting: add new candidate itemsets only when all of their subsets are estimated to be frequent

Articles about association rules and Python examples:

- <https://towardsdatascience.com/data-mining-market-basket-analysis-with-apriori-algorithm-970ff256a92c>
- <https://ml2021.medium.com/market-basket-analysis-apriori-algorithm-4fb052ffd2f2>
- <https://www.geeksforgeeks.org/apriori-algorithm/>
- <https://stackabuse.com/association-rule-mining-via-apriori-algorithm-in-python/>

Final remarks on Machine Learning

- All fields of ML are evolving quite fast at the moment as more companies want to utilize their data and big-tech is investing heavily
- Selection of some current trends in ML:
 1. Improved security: homomorphic encryption, federated learning, private set intersection, anonymizing data etc.
 2. Ethics and avoiding biases in AI
 3. Regulation of AI (e. g. on EU-level)
 4. Responsible/Sustainability in AI: use less computation/simpler models to reduce CO₂ emissions
 5. Few Shot, One Shot, & Zero Shot Machine Learning
 6. Understandable/Transparent AI
 7. Auto ML and democratizing AI
 8. Process adaptations like MLOps, DataOps, Data Governance
 9. AI Strategy on EU-, national- and company-level
 10. Tiny-ML, on-edge-inference/computing, specialized hardware (e. g. TPU)
 11. Multi-Modal Learning and Multi-Objective Models
 12. Quantum ML (hardware is not yet ready but algorithms are being adapted and software/middleware are being developed right now)

Final remarks on Machine Learning

- Some recommendations for continuous learning online:
 - <https://towardsdatascience.com/>
 - <https://www.kdnuggets.com/>
 - <https://www.kaggle.com/>
 - <https://www.geeksforgeeks.org/>
 - <https://www.reddit.com/r/datascience/>
 - <https://machinelearningmastery.com/>

and many more (also on youtube & Co)

Control Questions

You should be able to answer the following questions:

- How is unsupervised learning different from supervised and reinforcement learning?
- What is soft and hard clustering? Discuss a few use cases for clustering.
- How does k-means work and which variations do you know. Explain them and why they might make sense for a use case. What is the elbow method?
- How can you evaluate the quality of clustering results?
- What is hierarchical clustering? What two types exist? What is a dendrogram?
- How does DBSCAN work? What are its advantages?
- What are association rules? Where would you use association rules? Benefits and drawbacks?
- How does the apriori algorithm work and which principle does it use?

Links and Sources

- C. Bishop, Pattern Recognition and Machine Learning. 2006.
- Images and content adapted from slides of Nicole Artner, TU Wien, LV “Einführung Mustererkennung”
- https://en.wikipedia.org/wiki/K-means_clustering
- <https://en.wikipedia.org/wiki/K-means%2B%2B>
- [https://en.wikipedia.org/wiki/Silhouette_\(clustering\)](https://en.wikipedia.org/wiki/Silhouette_(clustering))
- https://en.wikipedia.org/wiki/Self-organizing_map
- Williams, J. T.; Carpenter, K. E.; Van Tassell, J. L.; Hoetjes, P.; Toller, W.; Etnoyer, P.; Smith, M. (2010). "Biodiversity Assessment of the Fishes of Saba Bank Atoll, Netherlands Antilles".
- <https://www.kdnuggets.com/2019/12/market-basket-analysis.html>
- https://en.wikipedia.org/wiki/Association_rule_learning
- <https://medium.com/analytics-vidhya/association-rule-mining-concept-and-implementation-28443d16f611>
- <https://www.kdnuggets.com/2016/04/association-rules-apriori-algorithm-tutorial.html/2>
- <https://www.youtube.com/watch?v=mtkWR8sx0NA>
- **Distances in higher dimensions and why they loose their “meaning”:**
<https://towardsdatascience.com/the-surprising-behaviour-of-distance-metrics-in-high-dimensions-c2cb72779ea6>